

ZigZag-Project: Enabling Fast Architecture-Mapping DSE for Deep Learning Accelerators

ISPASS2023 Tutorial

Arne Symons, Linyan Mei, Guilherme Paim

Prof. Marian Verhelst

MICAS Labs, ESAT, KU Leuven, Belgium



- Introduction (8:30 – 8:50)
- Lab 1: Assess HW performance of DNN layer onto accelerator, with fixed temporal mapping (8:50 – 9:30)
- Lab 2: Automate temporal mapping optimization (9:30 – 10:00)
- Break (10:00 to 10:30)
- Lab 3: Understand the HW architecture definition (10:30 – 11:00)
- Lab 4: Explore layer-fused mappings on multi-core architectures using Stream (11:00 – 11:45)
- Concluding remarks (11:45 – 12:00)

- Project started Dec. 2018
- Many contributors



Prof. Marian Verhelst



Pouya Houshmand



Steven Colleman



Vikram Jain




Guilherme Paim



Koen Goetschalckx



Arne Symons



Linyan Mei




Victor JUNG (ETHz)




Sebastian Karl (TUM)

Application




CAT, DOG, DUCK




Algorithm


Deep Feed Forward (DFF)




Deep Convolutional Network (DCN)




Auto Encoder (AE)






Long / Short Term Memory (LSTM)








Deconvolutional Network (DN)








Mapping/ Scheduling

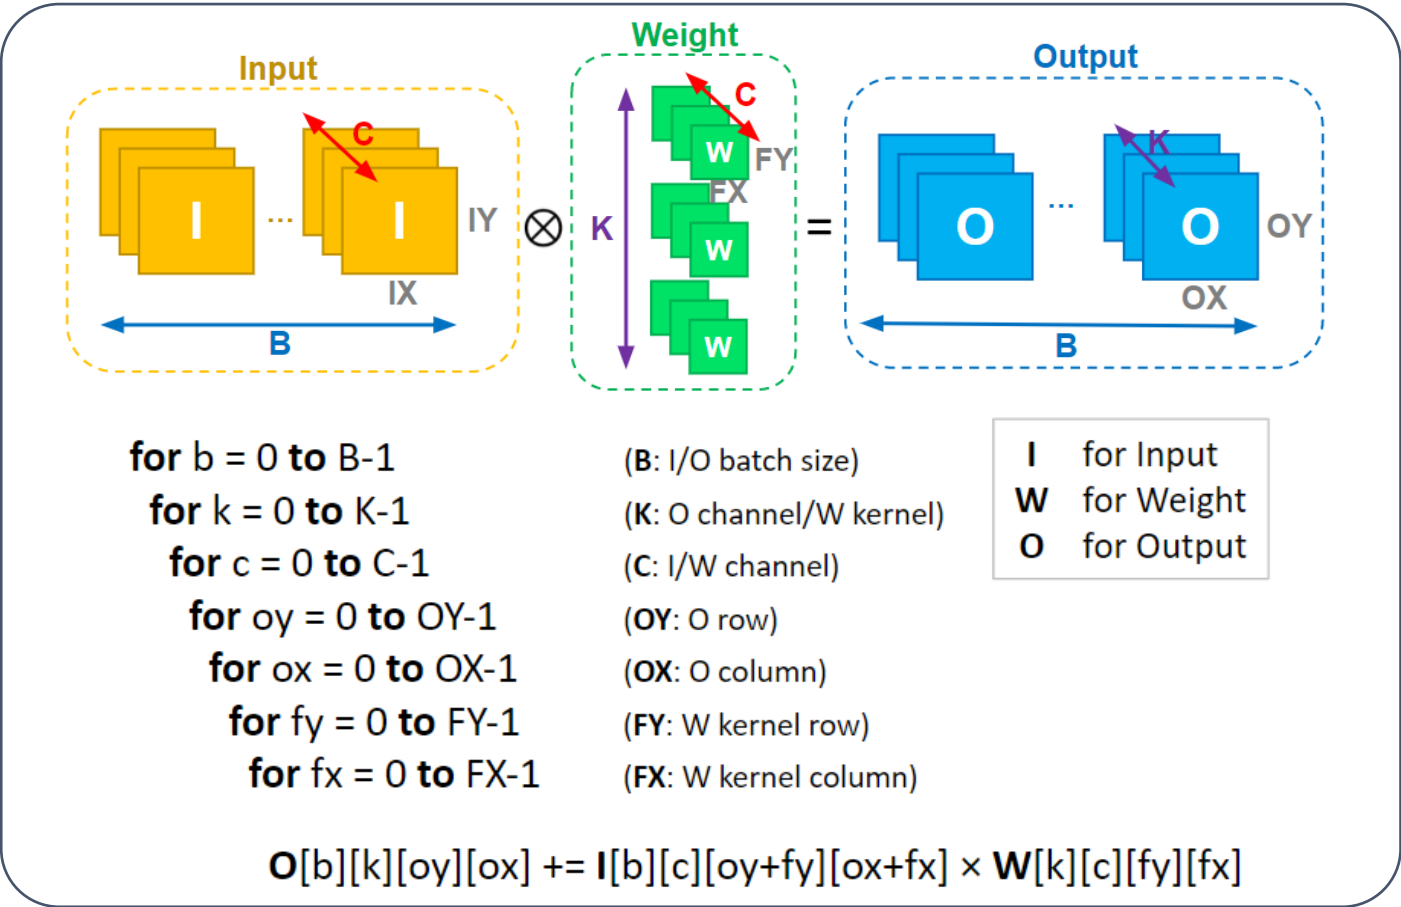


HW Arch.



Technology





	B	K	C	OY	OX	FY	FX
W	×	✓	✓	×	×	✓	✓
I	✓	×	✓	? ^{IY}	? ^{IX}	? ^{IY}	? ^{IX}
O	✓	✓	×	✓	✓	×	×

- ✓ relevant (r)
- × irrelevant (ir)
- ? partially relevant (pr)
- ?^{IX/IY} partially relevant to IX/IY

A DNN Conv2D layer:

3D operand (**W/I/O**) space.

7D nested for-loop
MAC operation.

Each Operand has its own
(ir)relevant loop dimensions.

- **r** loops contribute to **data size**.
- **ir** loops contribute to **data reuse**.
- **pr** loops contribute to both **data size** and **data reuse**.

Workload	I Batch size	O channel	I / W channel	O row	O column	W row	W column
Conv2D (right fig.)	B	K	C	OY	OX	FY	FX
Conv1D	B	K	C	1	OX	1	FX
Depthwise Conv2D*	B	1	1	OY	OX	FY	FX
Pointwise Conv2D	B	K	C	OY	OX	1	1
Matrix-Vector Multi.	1	K	C	1	1	1	1
Matrix-Matrix Multi.	B	K	C	1	1	1	1

* Repeat 'C' or 'K' times to finish one Depthwise Conv2D layer (C = K).

Conv2D

```
for b = 0 to B-1
  for k = 0 to K-1
    for c = 0 to C-1
      for oy = 0 to OY-1
        for ox = 0 to OX-1
          for fy = 0 to FY-1
            for fx = 0 to FX-1
              O[b][k][oy][ox] += I[b][c][oy+fy][ox+fx] x W[k][c][fy][fx]
```

(B: I/O batch size)
(K: O channel/W kernel)
(C: I/W channel)
(OY: O row)
(OX: O column)
(FY: W kernel row)
(FX: W kernel column)

I for Input
W for Weight
O for Output

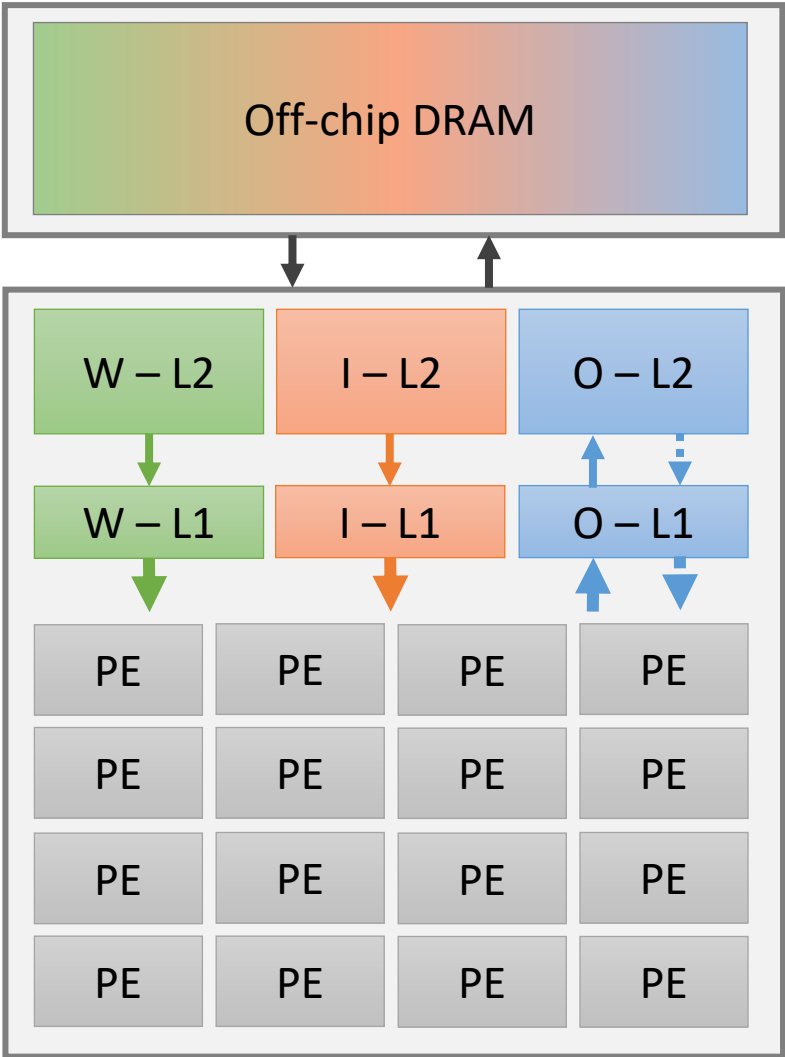
MMM

```
for b = 0 to B-1
  for k = 0 to K-1
    for c = 0 to C-1
      O[b][k] += I[b][c] x W[k][c]
```

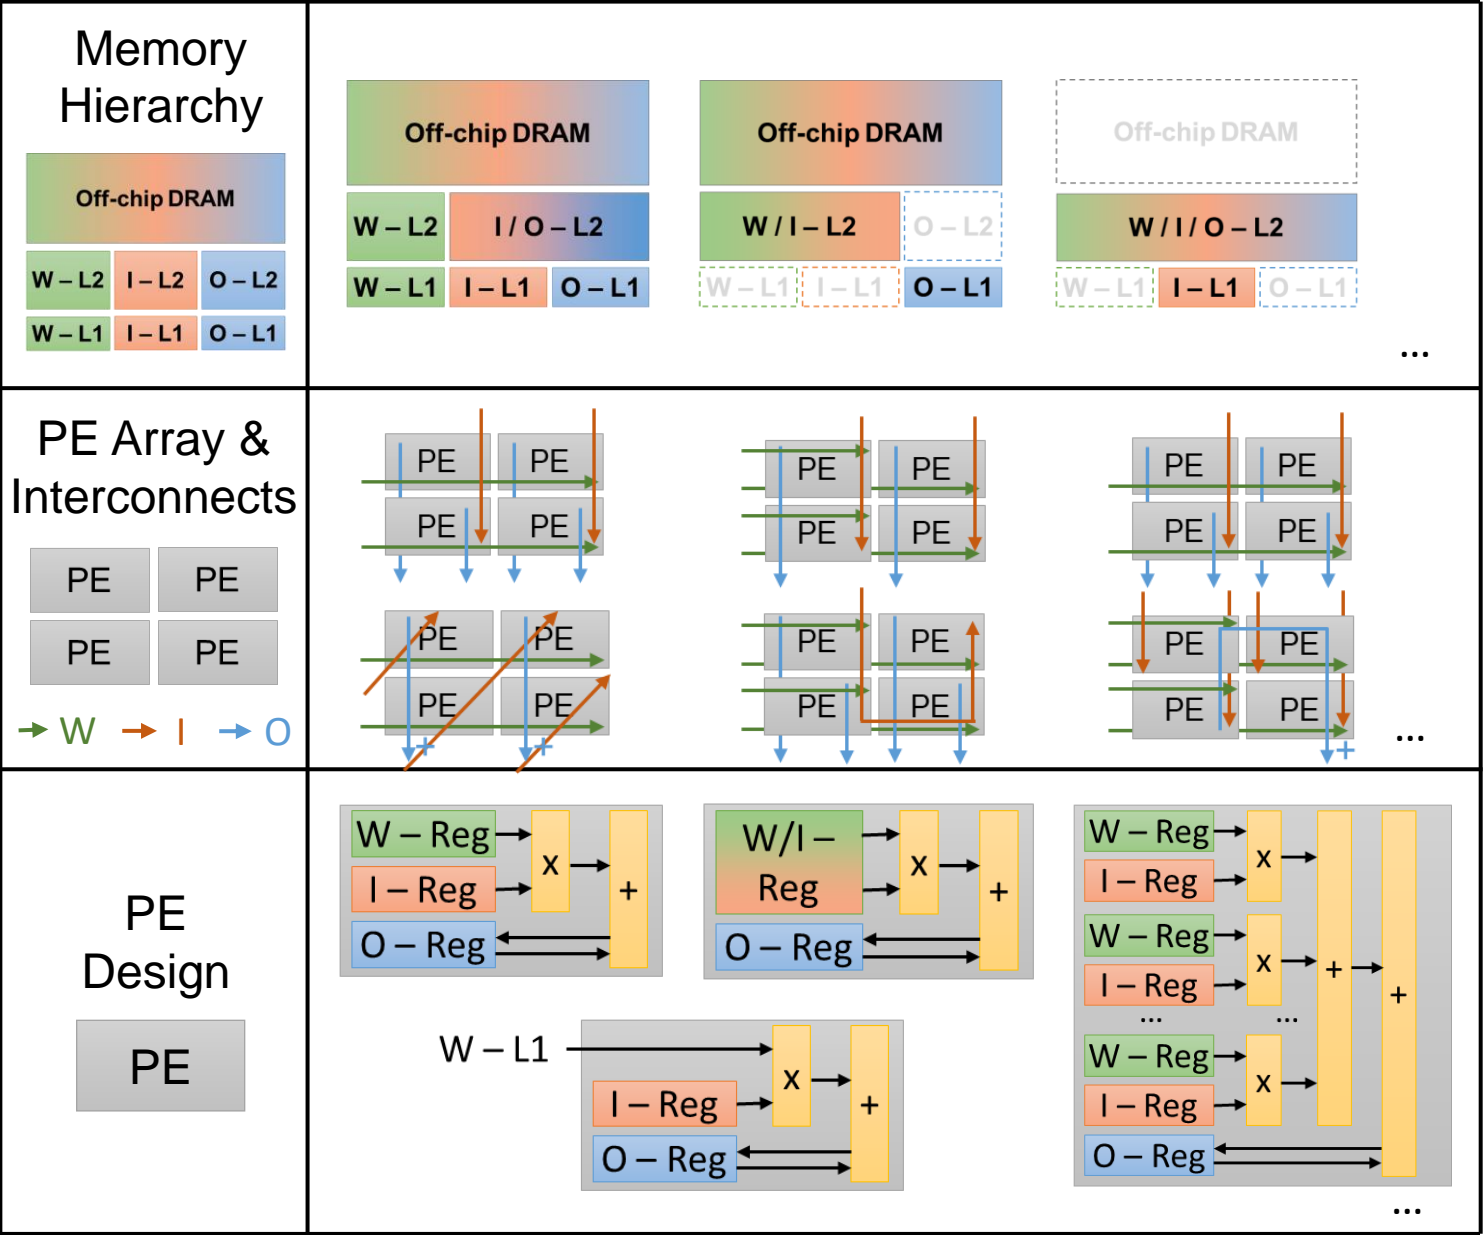
(B: I/O col)
(K: W/O row)
(C: W col / I row)

Most **ML workloads** fit into the regular **nested for-loop** format.

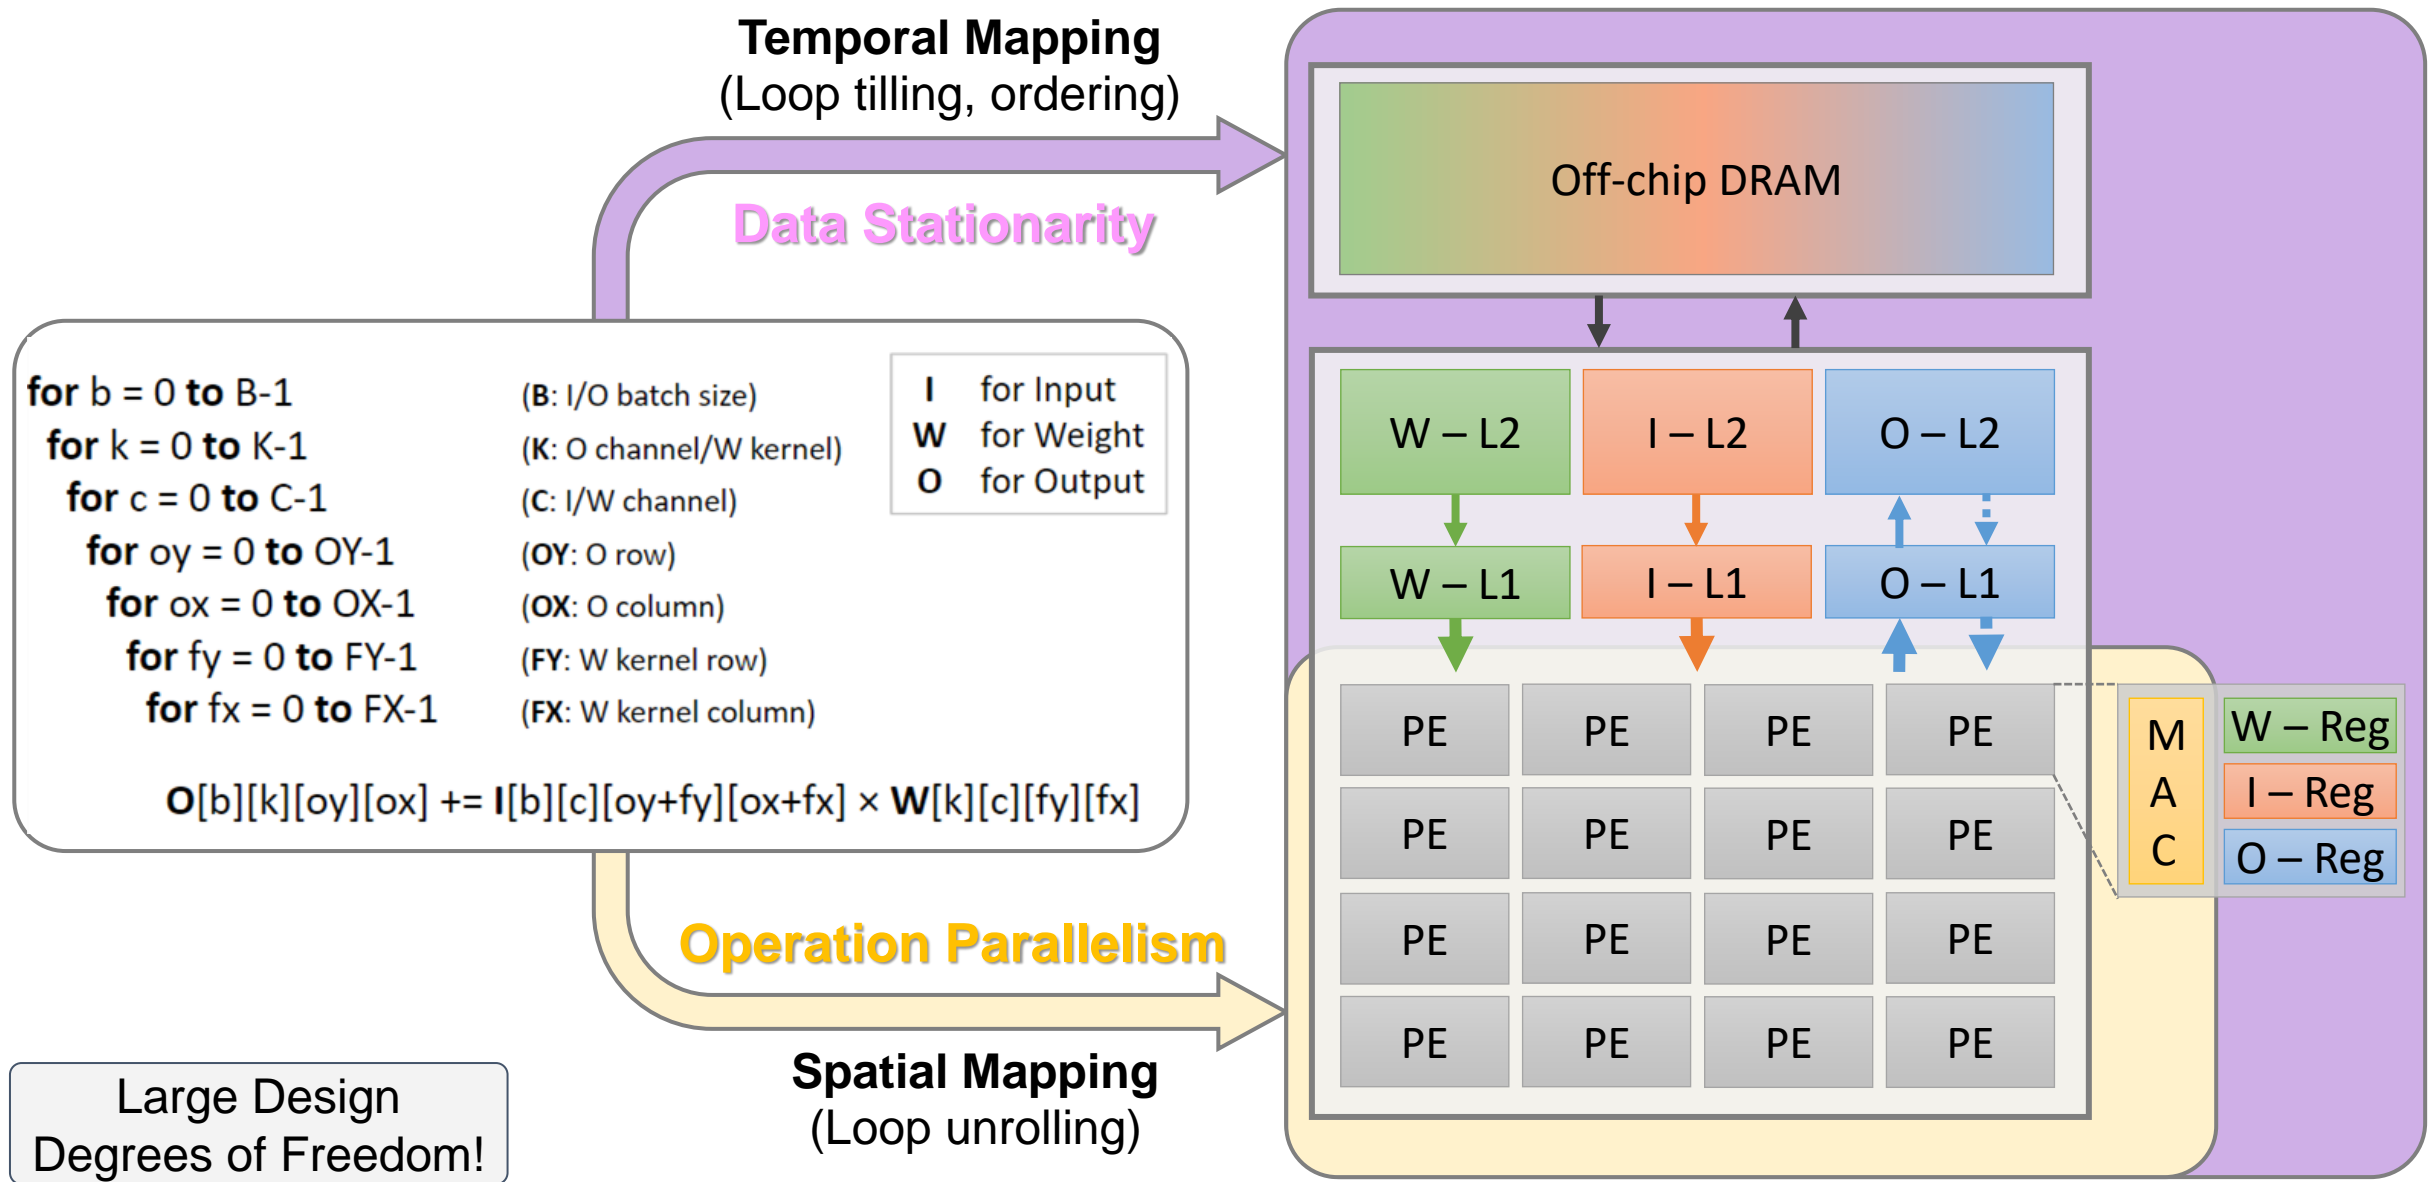
No data dependency between each for-loop.




Large Design Degrees of Freedom!




Mapping (a.k.a. Dataflow)







```
for b in [0: 3):  
  for k in [0: 4):  
    for c in [0: 8):  
      O[b][k] = I[b][c] × W[k][c]
```




```
for b in [0: 3):  
  for k in [0: 4):  
    for c in [0: 8):  
      O[b][k] = I[b][c] × W[k][c]
```



```
for b in [0: 3):  
  for c in [0: 8):  
    for k in [0: 4):  
      O[b][k] = I[b][c] × W[k][c]
```




```
for k in [0: 4):  
  for c in [0: 8):  
    for b in [0: 3):  
      O[b][k] = I[b][c] × W[k][c]
```

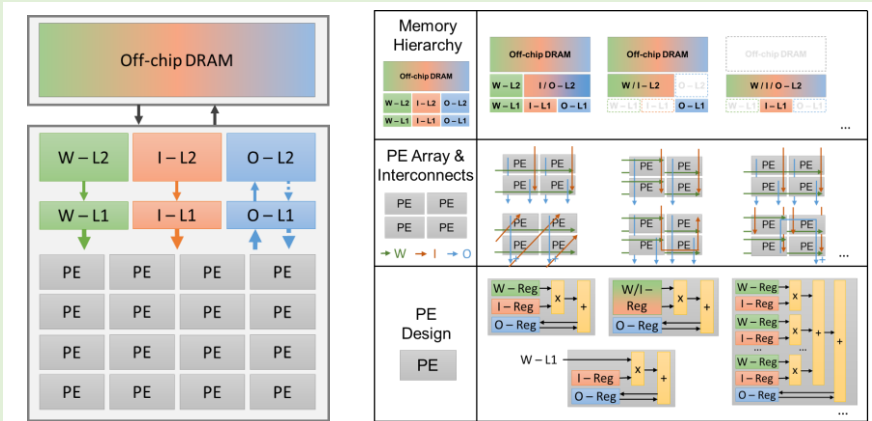


```
for k2 in [0: 2):  
  for c in [0: 8):  
    for b in [0: 3):  
      for k1 in [0: 2):  
        O[b][2k2+k1] = I[b][c] × W[2k2+k1][c]
```

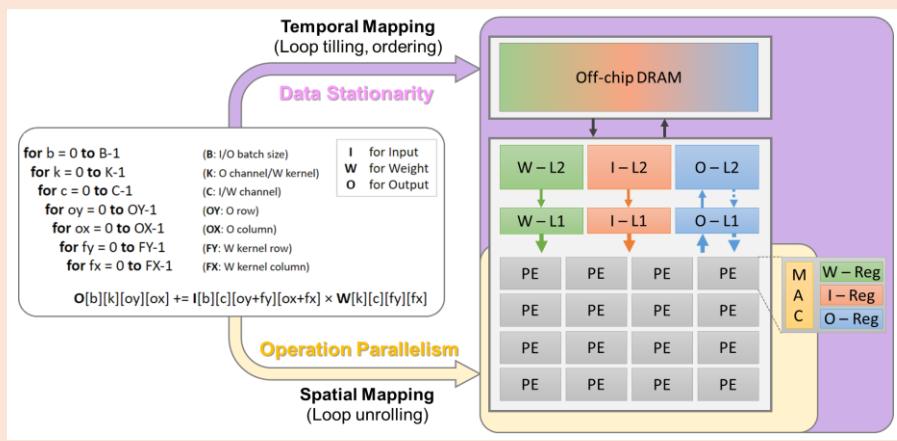
Algorithm



Hardware



Mapping



Technology and Others

Technology: 65nm/40nm/28nm/..., NVM, CIM, 3D IC, etc.

Others: Sparsity, various precisions, cross-layer execution, etc.

HUGE design space at each level & at combined levels.

Regular workload & Deterministic processing flow & Well-defined HW components.

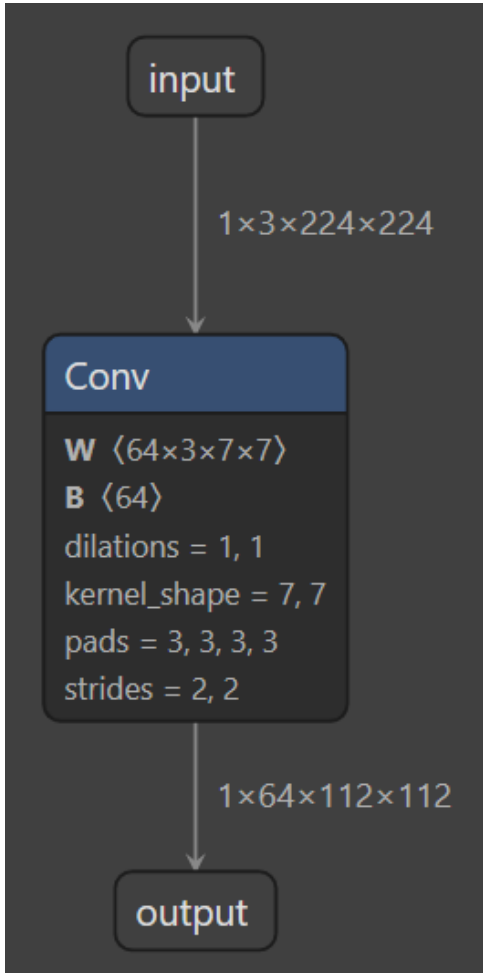
<https://github.com/ZigZag-Project/zigzag>

```
$ git clone git@github.com:ZigZag-Project/zigzag.git
$ cd zigzag
$ conda create --name my-zigzag-env python=3.10
$ conda activate my-zigzag-env
$ pip install -r requirements.txt
$ git checkout ispass2023-tutorial
$ code .
```

- Open lab1/main.py
- Expects three **arguments**:
 - accelerator
 - model
 - mapping
- Extracts names from the given arguments and sets inputs
- Defines the sequence of **stages** to be executed
- Runs the sequence of stages with inputs
- Plots the returned **CostModelEvaluation** (CME)

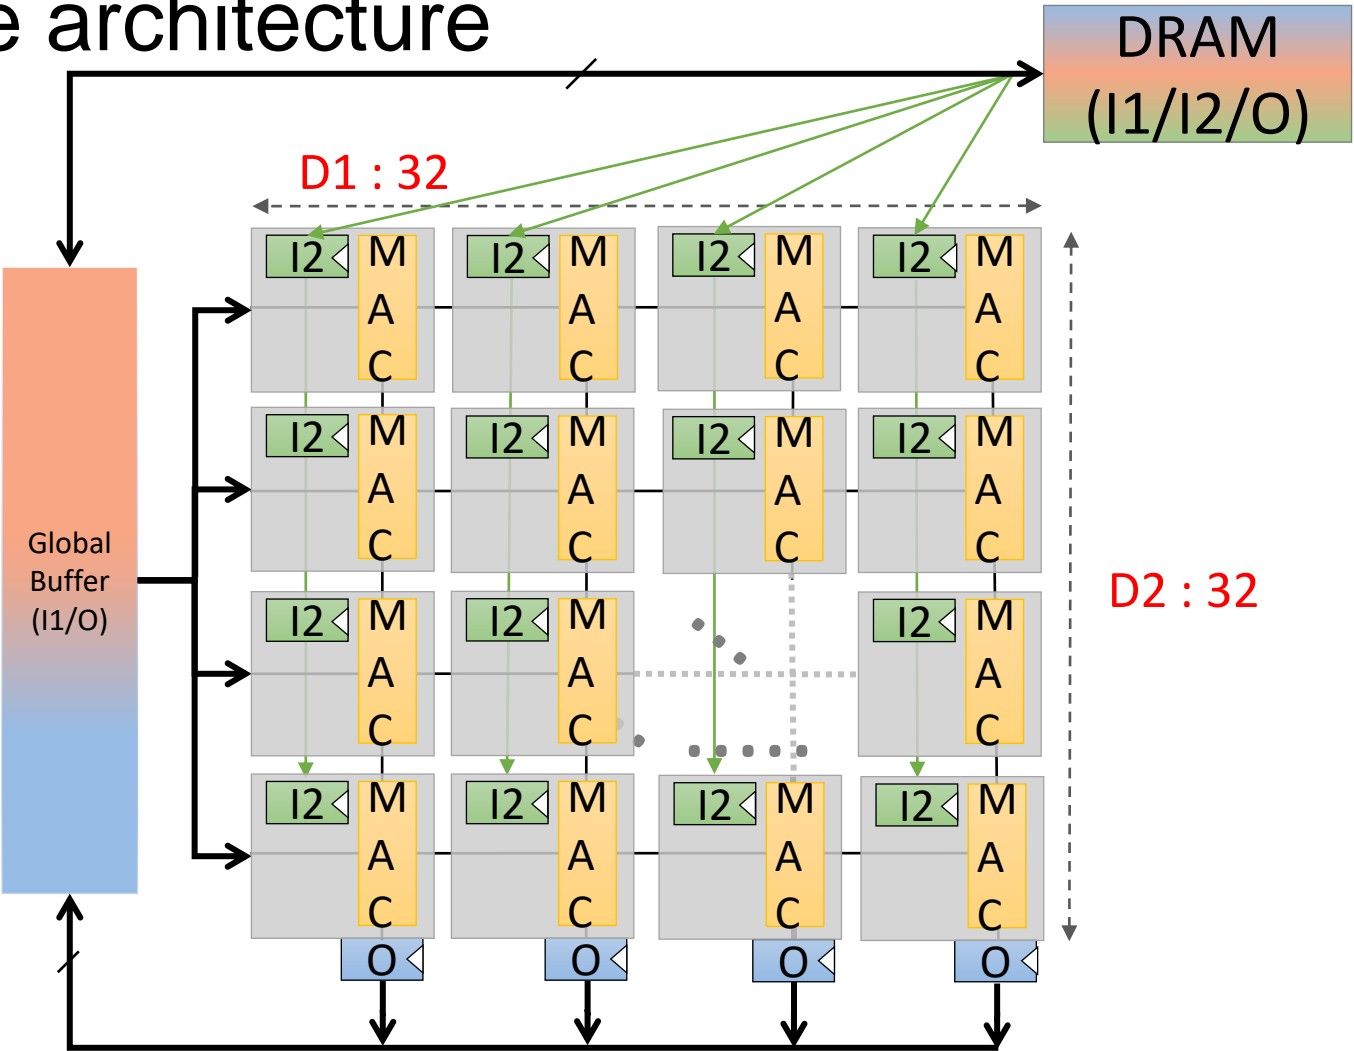
Model (workload)

- First layer of ResNet18 (ONNX format)



Accelerator

- TPU-like architecture



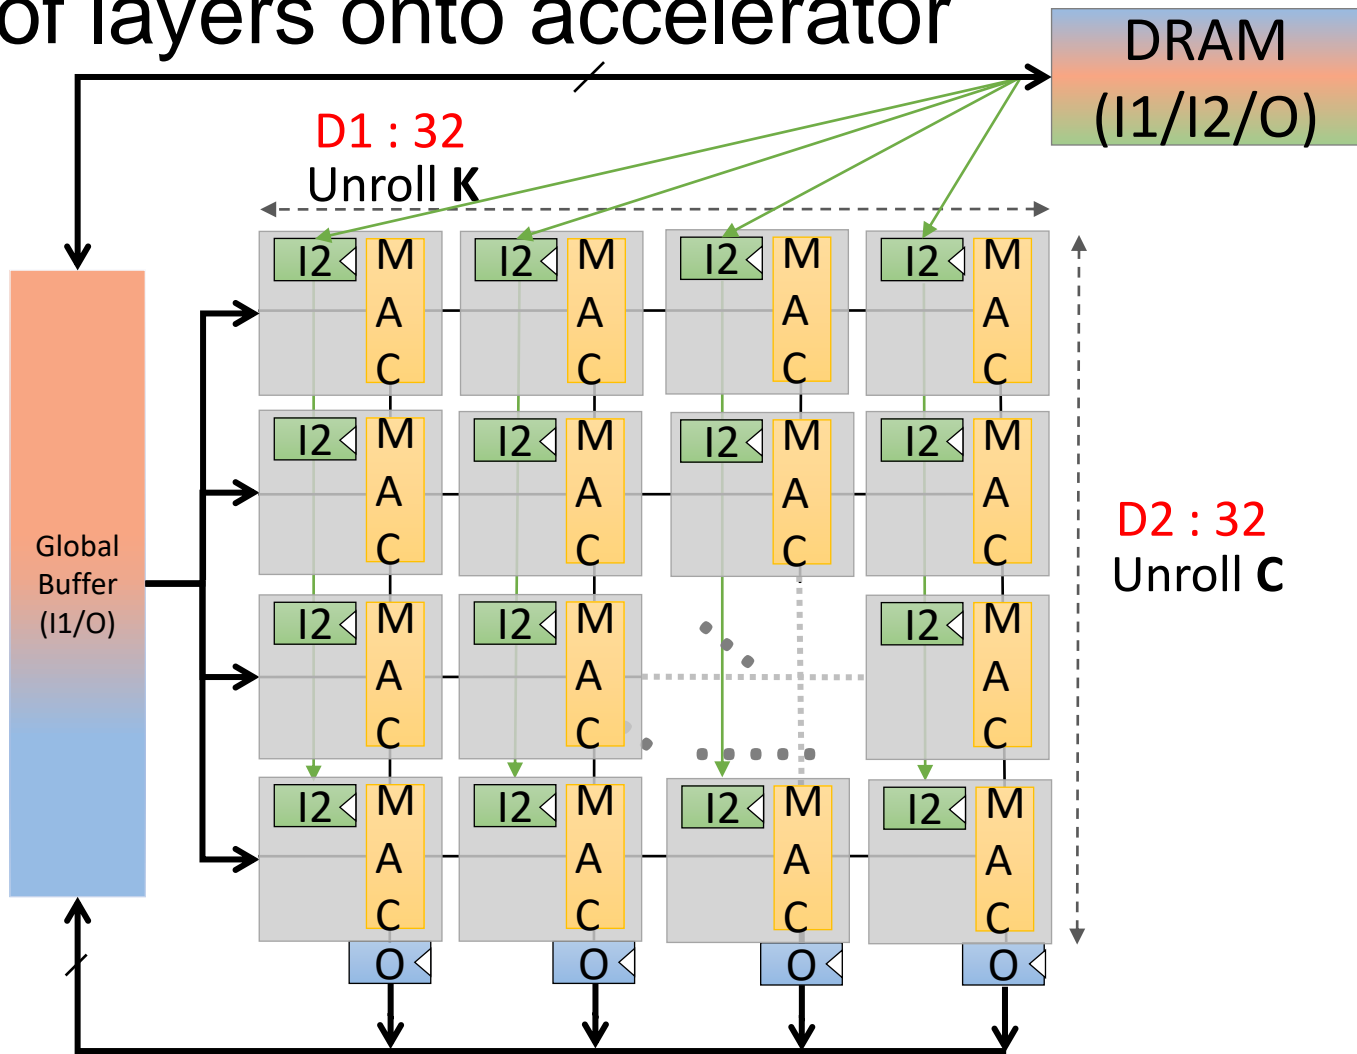
Mapping

- Defines mapping of layers onto accelerator

```
mapping = {  
    "/conv1/Conv": { # first ResNet18 layer name in onnx model  
        "spatial_mapping": {"D1": ("K", 32), "D2": ("C", 32)},  
        "temporal_ordering": [  
            # Innermost loop  
            ("OX", 112),  
            ("OY", 112),  
            ("FX", 7),  
            ("FY", 7),  
            ("K", 2),  
            # Outermost loop  
        ],  
    },  
}
```

Mapping

- Defines mapping of layers onto accelerator



First experiment:

- model = "lab1/resnet18_first_layer.onnx"
- accelerator = "zigzag.inputs.examples.hardware.TPU_like"
- mapping = "mapping"

- Run lab1/main.py

```
(my-zigzag-env) asymons@micaszb03:~/zigzag$ python lab1/main.py --model lab1/resnet18_first_layer.onnx --accelerator zigzag.inputs.examples.hardware.TPU_like --mapping mapping
```

Second experiment:

- Modify lab1/mapping.py:
 - Change temporal loop ordering
 - Run lab1/main.py

```
(my-zigzag-env) asymons@micaszb03:~/zigzag$ python lab1/main.py --model lab1/resnet18_first_layer.onnx --accelerator zigzag.inputs.examples.hardware.TPU_like --mapping mapping
```

- Copy lab1/main.py → lab2/main.py
- Replace TemporalOrderingConversionStage → LomaStage
- Change **dump_filename_pattern**
- Change plotting **save_path**

- Copy lab1/mapping.py → lab2/mapping.py
- Remove **temporal_ordering**

First experiment:

- model = "lab2/resnet18_first_layer.onnx"
- accelerator = "zigzag.inputs.examples.hardware.TPU_like"
- mapping = "mapping"

- Run lab2/main.py

```
(my-zigzag-env) asymons@micaszb03:~/zigzag$ python lab2/main.py --model lab2/resnet18_first_layer.onnx --accelerator zigzag.inputs.examples.hardware.TPU_like --mapping mapping
```


- Copy lab2/main.py → lab2/main_user_defined.py
- Replace ONNXModelParserStage → WorkloadParserStage
- Change **dump_filename_pattern**
- Change plotting **save_path**

Second experiment:

- model = “resnet18_first_layer”
- accelerator = “zigzag.inputs.examples.hardware.TPU_like”
- mapping = “mapping”

- Run lab2/main_user_defined.py

```
(my-zigzag-env) asymons@micaszb03:~/zigzag$ python lab2/main_user_defined.py --model resnet18_first_layer --accelerat  
or zigzag.inputs.examples.hardware.TPU_like --mapping mapping
```

- ZigZag is also distributed on PyPI

```
~/zigzag$ pip install zigzag-dse
```

- API call for common use-case

```
from zigzag.api import get_hardware_performance_zigzag
def get_hardware_performance_zigzag(
    ... workload,
    ... accelerator,
    ... mapping,
    ... opt="latency",
    ... dump_filename_pattern="outputs/{datetime}.json",
    ... pickle_filename="outputs/list_of_cmes.pickle",
):
```

- Lab 3 & 4 after the break
- Start at 10.30 AM

- Open lab3/inputs/hardware/...
 - Definition of multiplier array
 - Definition of memory hierarchy
 - Definition of core
- Open lab3/inputs/mapping/...
 - Definition of spatial mappings
- Open lab3/main.py
 - Uses API call for every core architecture mapping
 - Uses API call for architecture comparison plot

First experiment:

- Run lab3/main.py


```
(my-zigzag-env) asymons@micaszb03:~/zigzag$ python lab3/main.py
```


- Hardware accelerator model based on array of multipliers and attached memory hierarchy
- Hardware performance estimation of DNN layer through analytical cost model
- Optimization of layer mapping through different stages
- Enables co-exploration of accelerator & mapping




Frameworks	Workload	Hardware	Mapping
ZigZag	A NN layer	Single-core accelerator	Single-layer mapping
Stream	One/more NN(s)	Multi-core accelerator	Fine-grained layer-fused mapping

Focus on **Stream** for the rest of the session


An example workload




An example accelerator






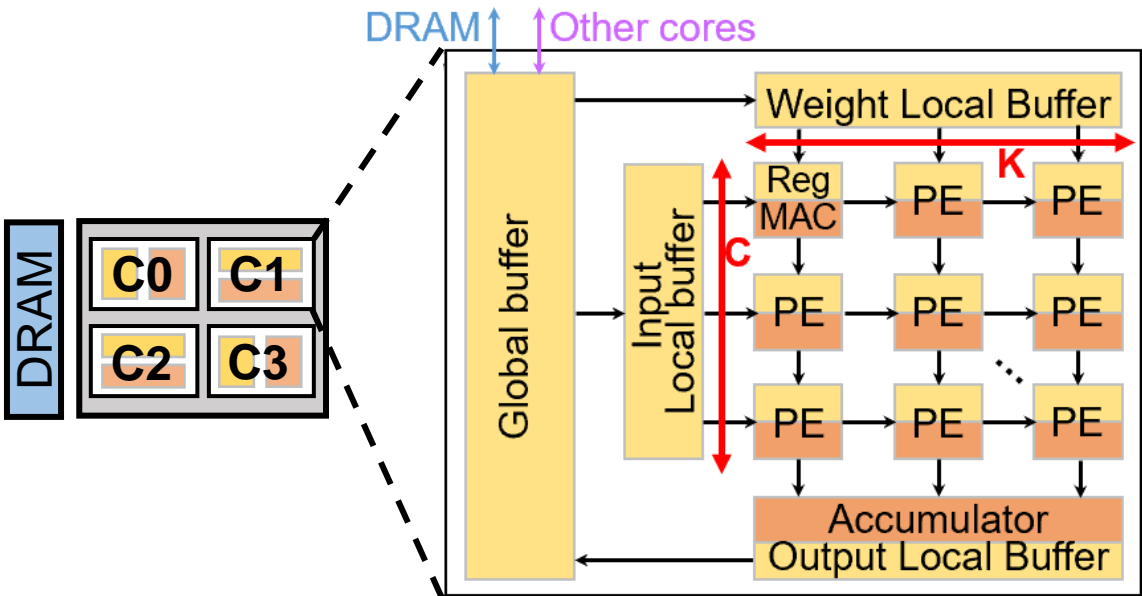
Acceleration timeline




An example workload





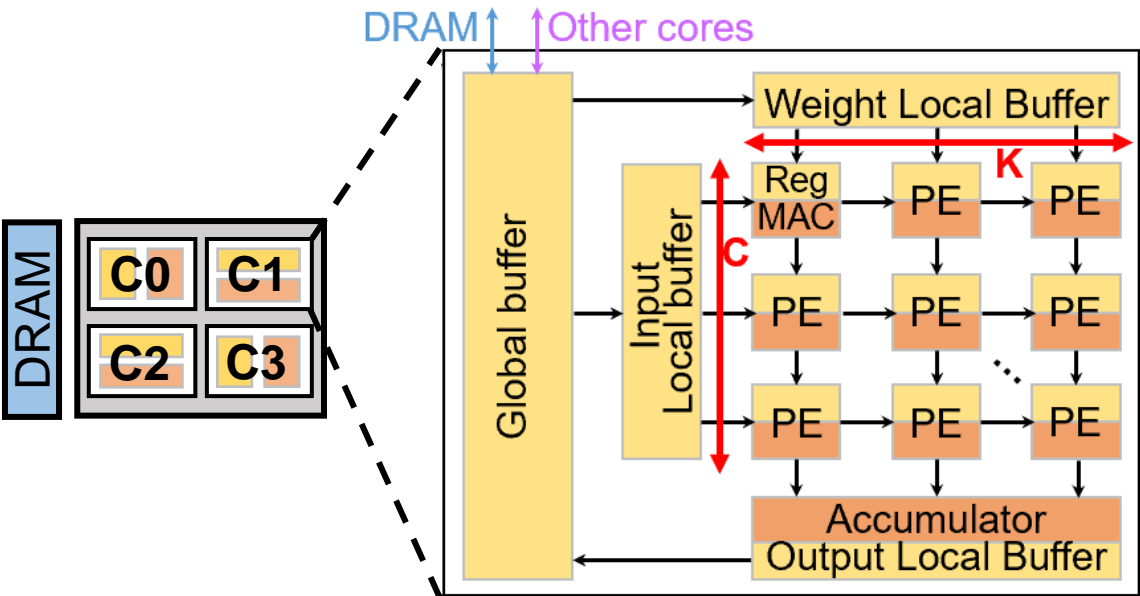
A multi-core accelerator




Tiled for layer-fused processing




A multi-core accelerator




Acceleration timeline



An example workload




Tiled for layer-fused scheduling



Schedule the workload to hardware accelerators

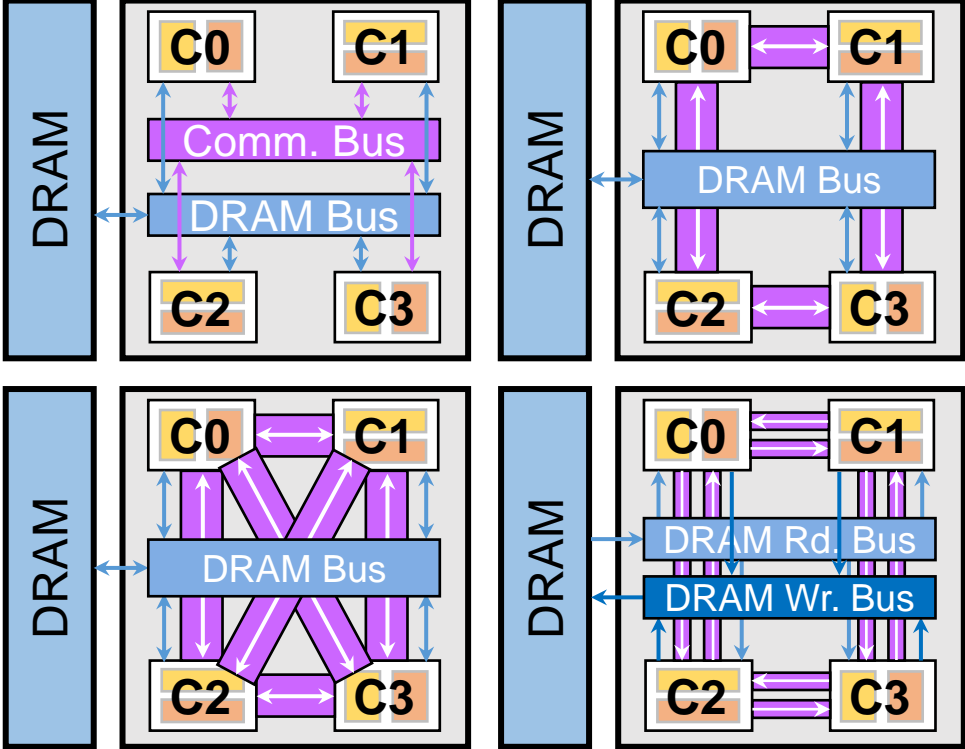
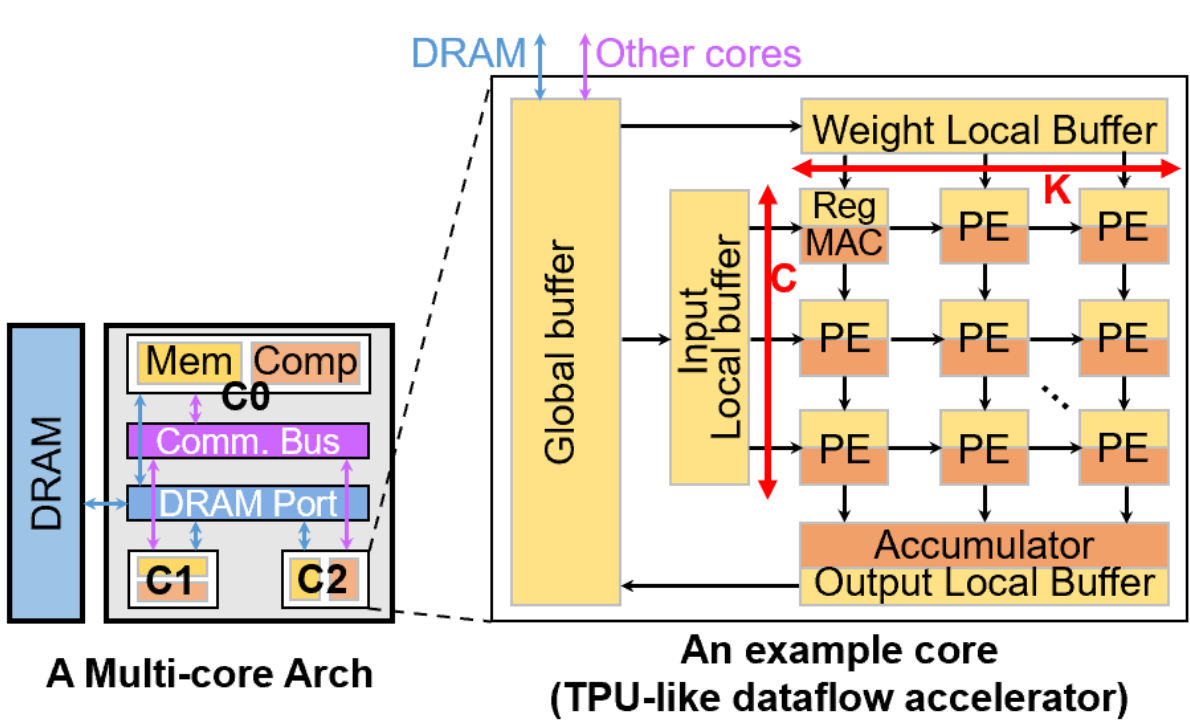
Hardware Schedule	Single-core	Multi-core
Layer-by-layer	(1)	(3)
SotA frameworks	ZigZag, Timeloop, ..., Stream	Kwon et al. 2021, Stream
Layer-fused	(2)	(4)
SotA frameworks	DeFiNES, TVM-Cascade, Stream	Stream

- ✓ Model single-core architecture (identical to ZigZag)




What can Stream do?



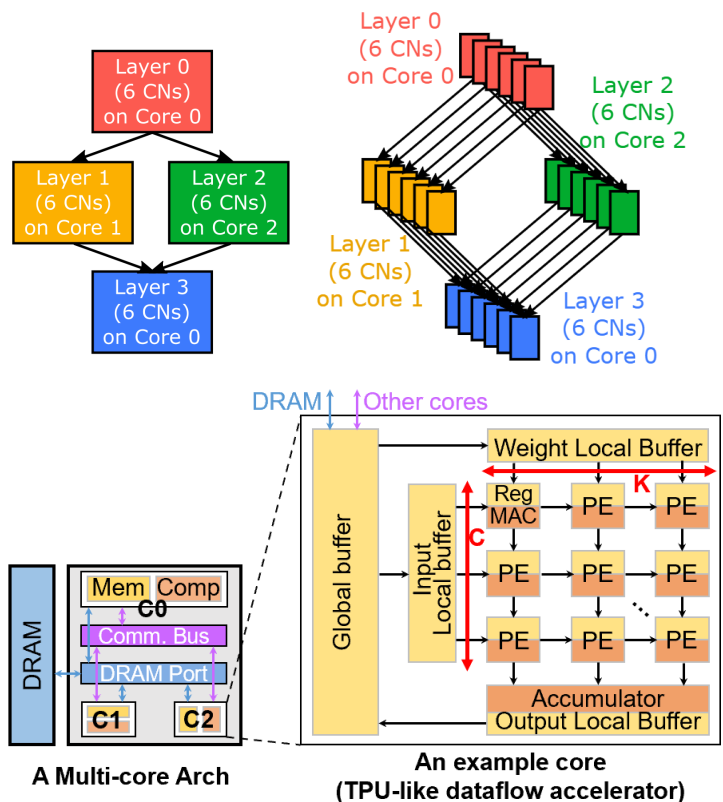
- ✓ Model single-core architecture (identical to ZigZag)
- ✓ Model different multi-core topologies




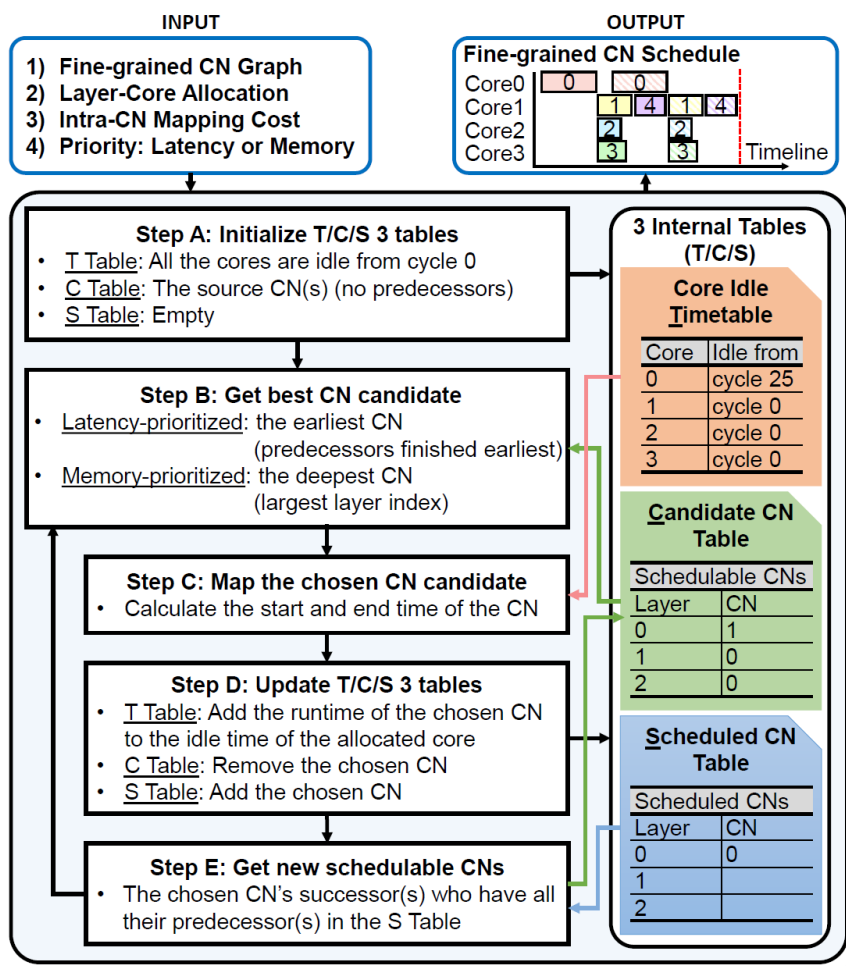
- ✓ Model single-core architecture (identical to ZigZag)
- ✓ Model different multi-core topologies
- ✓ Model different scheduling granularities



- ✓ Model single-core architecture (identical to ZigZag)
- ✓ Model different multi-core topologies
- ✓ Model different scheduling granularities
- ✓ Model different scheduling heuristics



Stream Overview



<https://github.com/ZigZag-Project/stream>

```
$ git clone git@github.com:ZigZag-Project/stream.git
$ cd stream
$ conda create --name my-stream-env python=3.10
$ conda activate my-zigzag-env
$ pip install -r requirements.txt
$ git checkout ispass2023-tutorial
$ code .
```

- Open lab4/main_fixed.py
- Defines inputs directly in file instead of arguments
- Extracts (from input names) and defines variables
- Sets up the sequence of stages
- Runs the stages
- Plots the results

Workload

- Open
lab4/inputs/workload/duplicated_resnet18_layer_fixed.py
- First layer of ResNet18 duplicated 4 times
- Dependencies between the layers (**operand_source**)
- **operator_type** overloaded for fixed mapping

Accelerator

- Open `lab4/inputs/hardware/heterogeneous_quadcore.py`
- Imports the “computational cores”
- Imports the pooling and simd cores
- Imports the offchip core
- Creates a 2D mesh of these cores
- Defines the multi-core Accelerator object

Mapping

- Open lab4/inputs/mapping/mapping_fixed.py
- Defines for each **operator_type** the possible layer-core allocations

First experiment:

- model = "...duplicated_resnet18_layer_fixed"
- accelerator = "...heterogeneous_quadcore"
- mapping = "...mapping_fixed"
- Run lab4/main_fixed.py

```
(my-stream-env) asymons@micaszb03:~/stream$ python lab4/main_fixed.py
```

Second experiment:

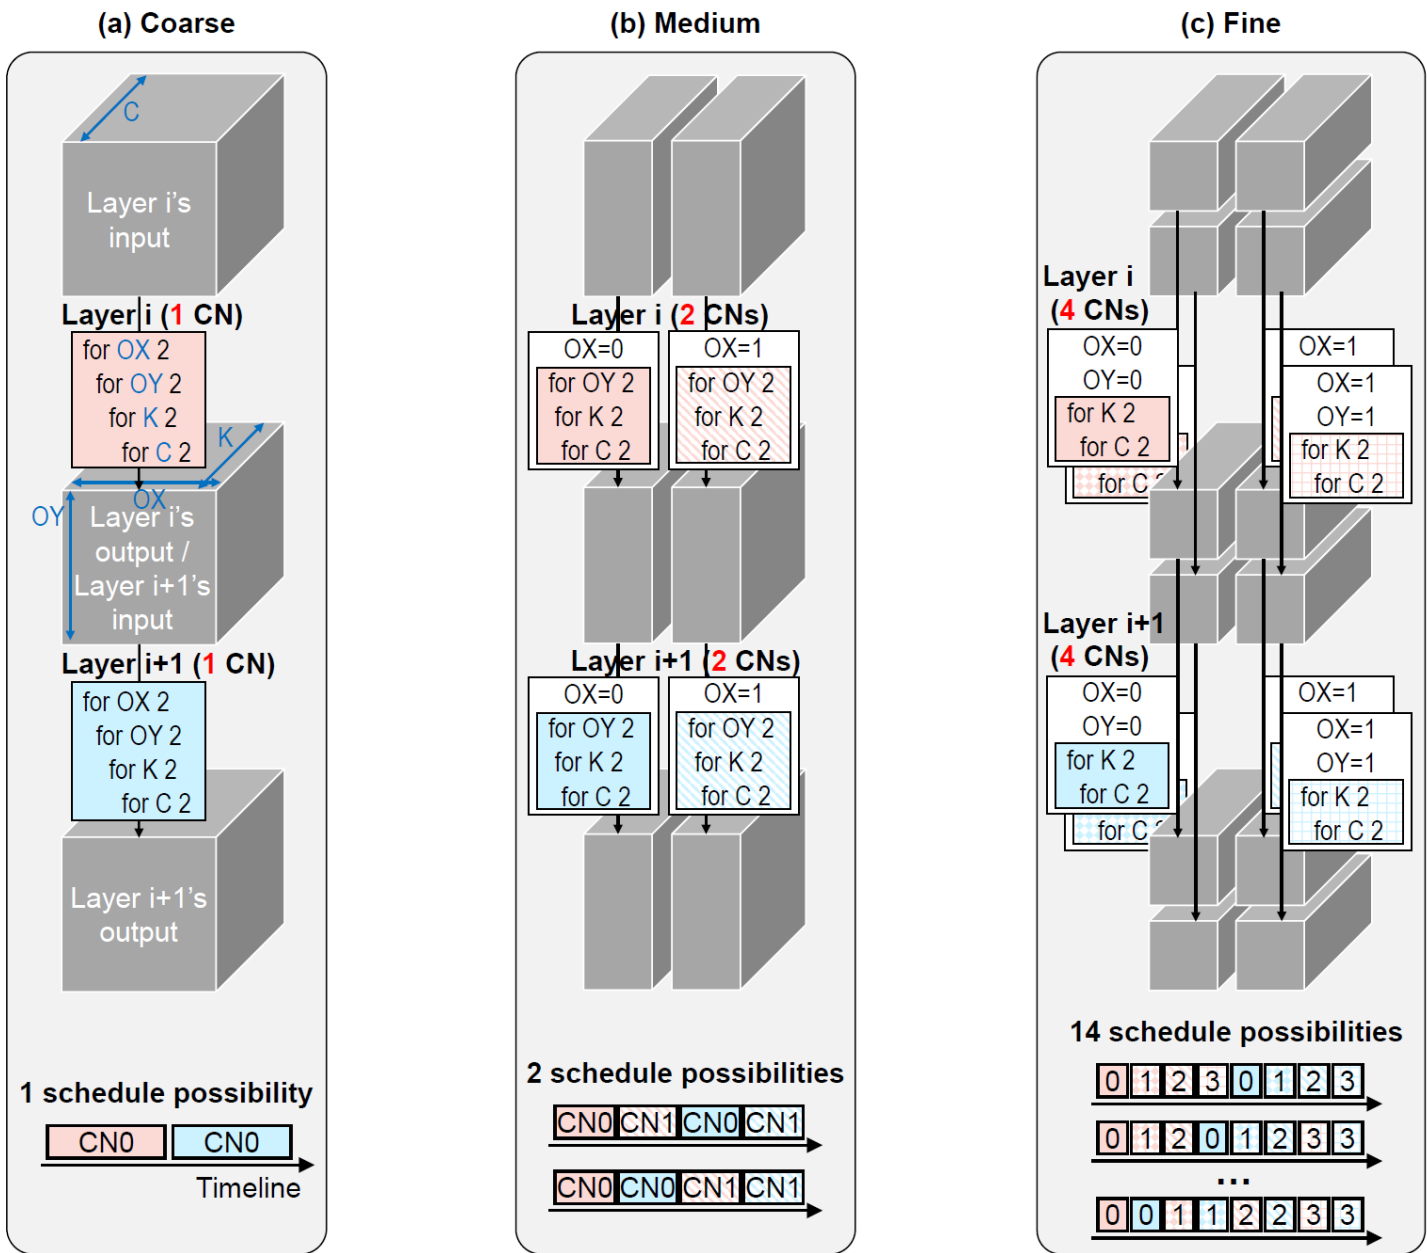
- What happens if we remove the layer dependencies?
- Remove the **operand_source** and modify the **constant_operands**
- Run lab4/main_fixed.py

```
(my-stream-env) asymons@micaszb03:~/stream$ python lab4/main_fixed.py
```

Third experiment:

- Allow genetic algorithm to find best layer-core allocation
- model = "...duplicated_resnet18_layer"
 - Modified **operator_type**
- mapping = "...mapping"
 - Modified for flexible layer-core allocation
- Run lab4/main_layer_by_layer.py

```
(my-stream-env) asymons@micaszb03:~/stream$ python lab4/main_layer_by_layer.py
```



Computation node (CN) granularity impacts scheduling flexibility

and others...

- Core utilization
- Intra-CN data reuse
- Data loading overhead
- Control overhead

- Open lab4/main_layer_fused.py
- **hint_loops** defines the outer-CN loops
 - `hint_loops = [("OY", 2)]` means 2 CNs per layer
 - `hint_loops = [("OY", "all")]` means OY CNs per layer

Fourth experiment:

- Assess the impact of layer-fused processing
- Run lab4/main_layer_fused.py

```
(my-stream-env) asymons@micaszb03:~/stream$ python lab4/main_layer_fused.py
```


- Open lab4/main.py
- End-to-end ResNet18 onnx model
- Layer-by-layer

Last experiment:

- End-to-end ResNet18 example
- Run lab4/main.py (layer-by-layer)
- If time permits:
 - Modify hint_loops with (“OY”, “all”)
 - Re-run and analyze differences

- Automatically infer optimal CN granularity
- Integrate inter-core connect energy estimation framework
- Automatically search for optimal multi-core architectures
- Add optimization constraints (e.g. max latency, area, ...)

- Code generation for existing accelerators
- Automatic hardware generation from hardware templates

- ZigZag enables fast hardware performance estimation for specialized DNN accelerator architectures
- Mapping optimizations yield better performance
- Co-exploration of architecture with mapping

- Stream extends the capabilities to multi-core architectures employing layer-fused scheduling
- Unified hardware model for different architecture topologies
- Different scheduling granularities through computation node

- **Goal: Enabling Fast Architecture-Scheduling/Mapping DSE for Machine Learning Accelerators**
- Github: <https://github.com/ZigZag-Project>
 - ZigZag
 - ZigZag-Demo
 - DeFiNES
 - Stream
- ZigZag documentation: <https://zigzag-project.github.io/zigzag/>
- Stream documentation: Underway (end of May)
- ZigZag-related publications:
<https://zigzag-project.github.io/zigzag/publications.html>